

COMP2611: Computer Organization

Introduction to the Pipelined Processor

- ❑ Two techniques for designing high-performance processors:
 - ❑ **Pipelining**
 - ❑ **Multiprocessing**

- ❑ Both techniques exploit **parallelism**:
 - ❑ **Pipelining**: parallelism among multiple instructions
 - ❑ **Multiprocessing**: parallelism among multiple processors

- ❑ In this course, we only focus on **pipelining**






- ❑ Single Cycle Datapath: simple but inefficient. Longest path (consists of 5 execution stages) determines the duration of the clock cycle
 - ❑ Not every instruction needs all the 5 stages, yet every instruction takes a complete clock cycle => **unnecessary performance penalty!**
- ❑ Multi-cycle Datapath: more efficient, but still wastes resources
 - ❑ **Functional units may remain idle in some clock cycles** until the completion of the Instruction (e.g., ALU unused in stages 4 and 5, the register file is unused in stages 3, 4, 5)

Pipelining can improve the performance further, in terms of throughput and efficiency by reusing functional units for other instructions


1- Introduction to pipelining

- ❑ Let's say we have 4 batches of laundry work

In serial



Pipeline




Startup time: time needed to fill the pipeline

Key characteristics:


- ❑ **Multiple tasks** are processed simultaneously
- ❑ Ideally, these tasks should be **independent** of each other
- ❑ Pipelining **does not help the latency** of a single task
- ❑ But, it **helps the throughput** of the entire workload
- ❑ Completion order in pipelined execution = that in sequential execution

How much can a pipeline improve?

- ❑ **Potential speedup = number of pipeline stages**
- ❑ The pipeline rate is limited by the slowest pipeline stage
 - Unbalanced lengths of pipeline stages can reduce speedup. Why?



- ❑ Can I align the pipeline stages like the above?
- ❑ Answer: **NO** because the tasks executing in parallel are not independent (task 4 overlaps task 4)
- ❑ The condition to align is to make sure **NO OVERLAP** of any stages?





❑ Pipeline performance (example):

Assume we require:

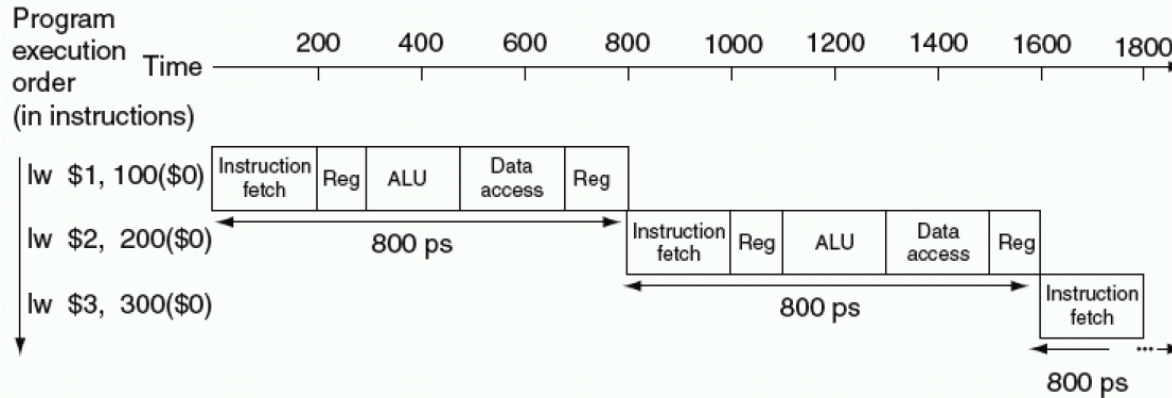
- 100 picoseconds for register read or write
- 200 picoseconds for all other stages

Instruction	Instruction fetch	Register read	ALU op	Memory access	Register write	Total time
Load Word (lw)	200ps	100 ps	200ps	200ps	100 ps	800ps
Store Word (sw)	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
Branch (beq)	200ps	100 ps	200ps			500ps

- ❑ The instruction **delays** in the example:
 - 800 ps (single cycle datapath)
 - 1000 ps (pipelined datapath)
- ❑ The instruction **throughputs** in the example:
 - 1 instruction per 800 ps (single cycle datapath)
 - 1 instruction per 200 ps (long time average for pipelined datapath)
- ❑ Pipelining does not improve the latency of a single instruction, it improves the throughput of the system (i.e., the datapath)
- ❑ In general (ideally), if we have a N stage pipeline:
 - ❑ We need N-1 cycles to fill the pipeline,
 - ❑ Then one instruction will finish per cycle
 - ❑ So the throughput is: $\text{Clock Rate} \times \text{IC} / (\text{IC} + N - 1)$, with $\text{IC} \gg N$



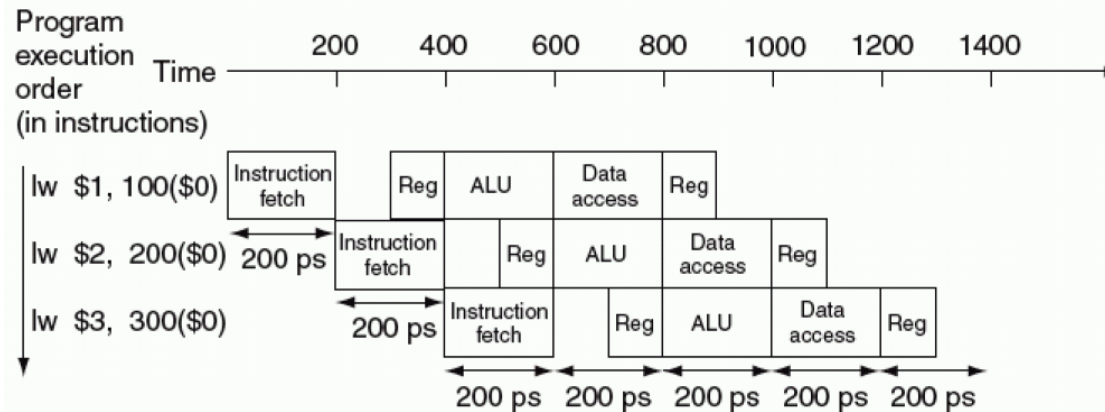
❑ Single-cycle, non-pipelined execution:



Total execution time
= 2400 ps

❑ Multicycle: > 2400 ps

❑ Pipelined execution



Total execution time
= 1400 ps

A program, 1000 instructions, all independent of each other
30% take 30ns, 40% take 40ns, 30% take 50ns

Single-cycle implementation, (cycle time = 50ns)

❑ $CPI = 1.0$,

❑ $Time = \# \text{ of instr.} * CPI * \text{cycle time} = 1000 * 1.0 * 50ns = 50ms$

5-stage multi-cycle implementation, (cycle time = 10ns)

❑ $CPI = 0.3 * 3 + 0.4 * 4 + 0.3 * 5 = 4.0$

❑ $Time = 1000 * 4.0 * 10ns = 40ms$

5-stage pipeline implementation, (cycle time = 10ns)

❑ $CPI = (4 \text{ cycles } /*\text{startup}*/ + 1000 \text{ cycles } /*\text{pipelined}*/) / 1000$
 $= 1.004$

❑ $Time = 1000 * 1.004 * 10ns = 10.04ms$



ISA design affects the complexity of pipeline implementation.

MIPS ISA is designed for pipelining

- ❑ **All instruction are of the same length** (32-bit)

Easy to fetch one instruction in first stage of the pipeline and decode it in the second

- ❑ **It has just a few similar instruction formats**

With the source register fields being located in the same place in all instructions, 2nd stage can read the register file while decoding the type of instruction just fetched

- ❑ **Memory operands only appear in loads and stores**

We can use the execute stage to calculate the memory address and then access memory in the following stage

- ❑ **Alignment of memory operands on word boundaries**

We need not worry about a single data transfer instruction requiring two memory accesses; the data can be transferred between processor and memory in a single pipeline stage



Execution of each instruction is broken into 5 stages: (in the order of execution)

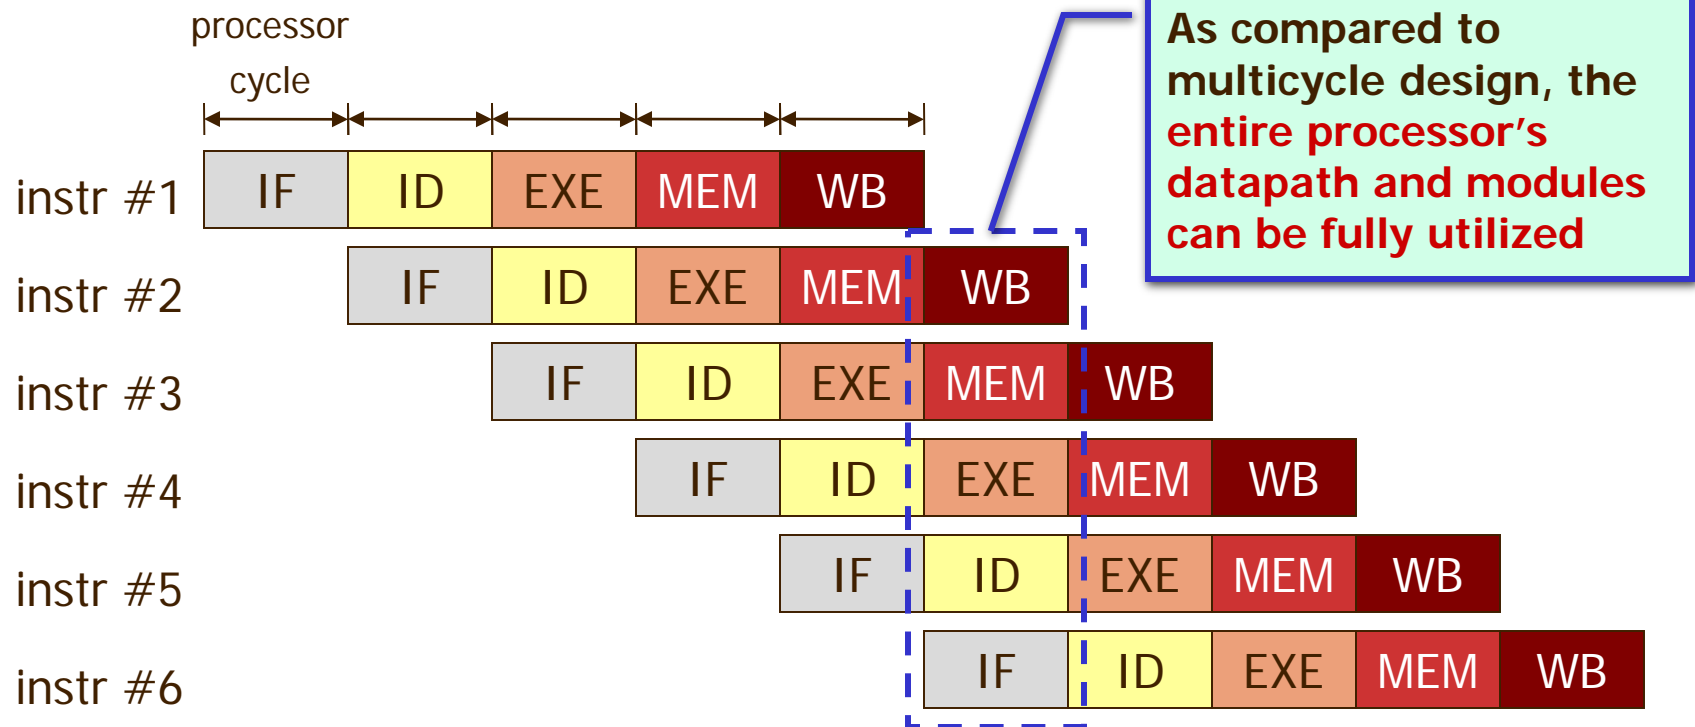
- **IF** : Fetch the instruction from memory
- **ID** : Instruction decode & register read
- **EX** : Perform ALU operation
- **MEM** : Memory access (if necessary)
- **WB** : Write result back to register

Each stage uses a different hardware unit and takes one clock cycle to complete.

- ❑ Instructions are allowed to share the different hardware units of the datapath as long as they are using the different hardware units (pipelining).

In pipelined processor,

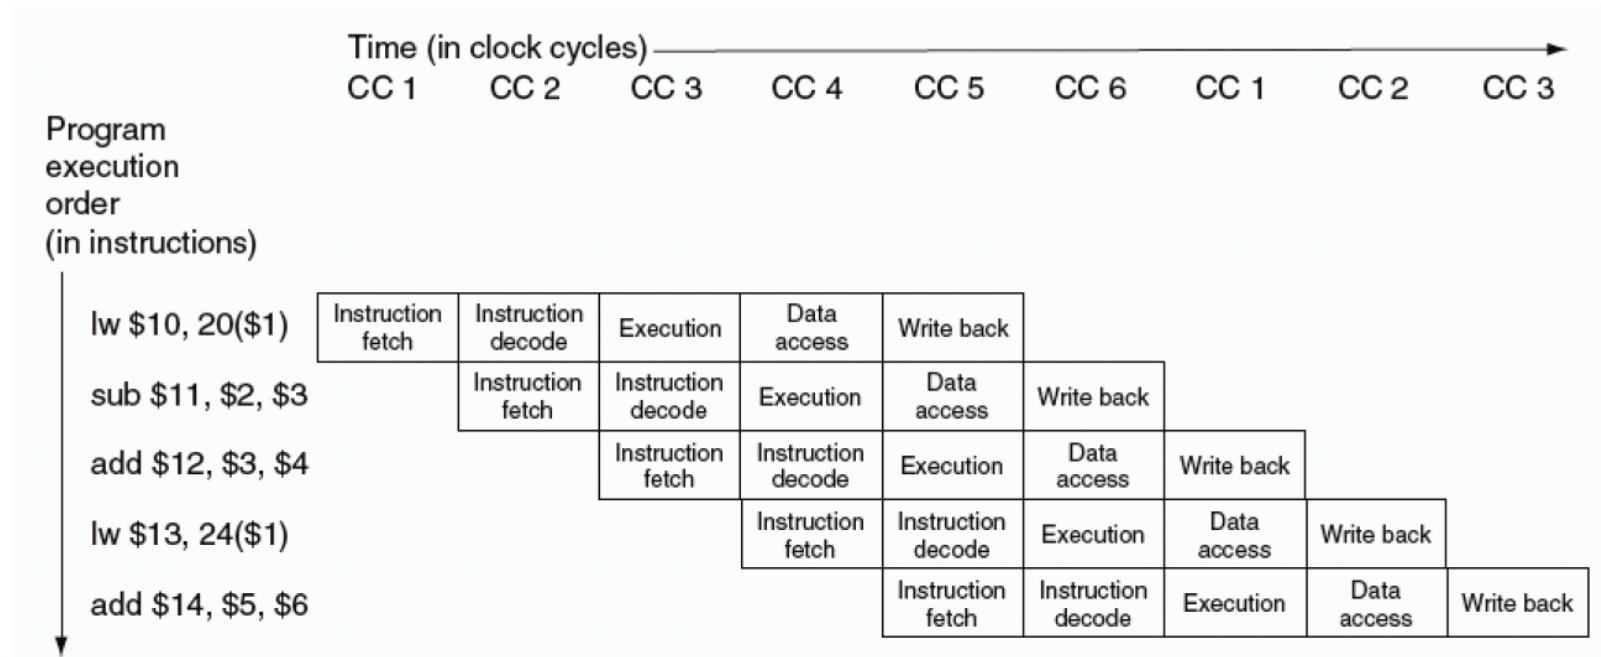
- ❑ Each instruction takes multiple steps (like **multicycle approach**)
- ❑ Each step is independent of each other and takes different datapath



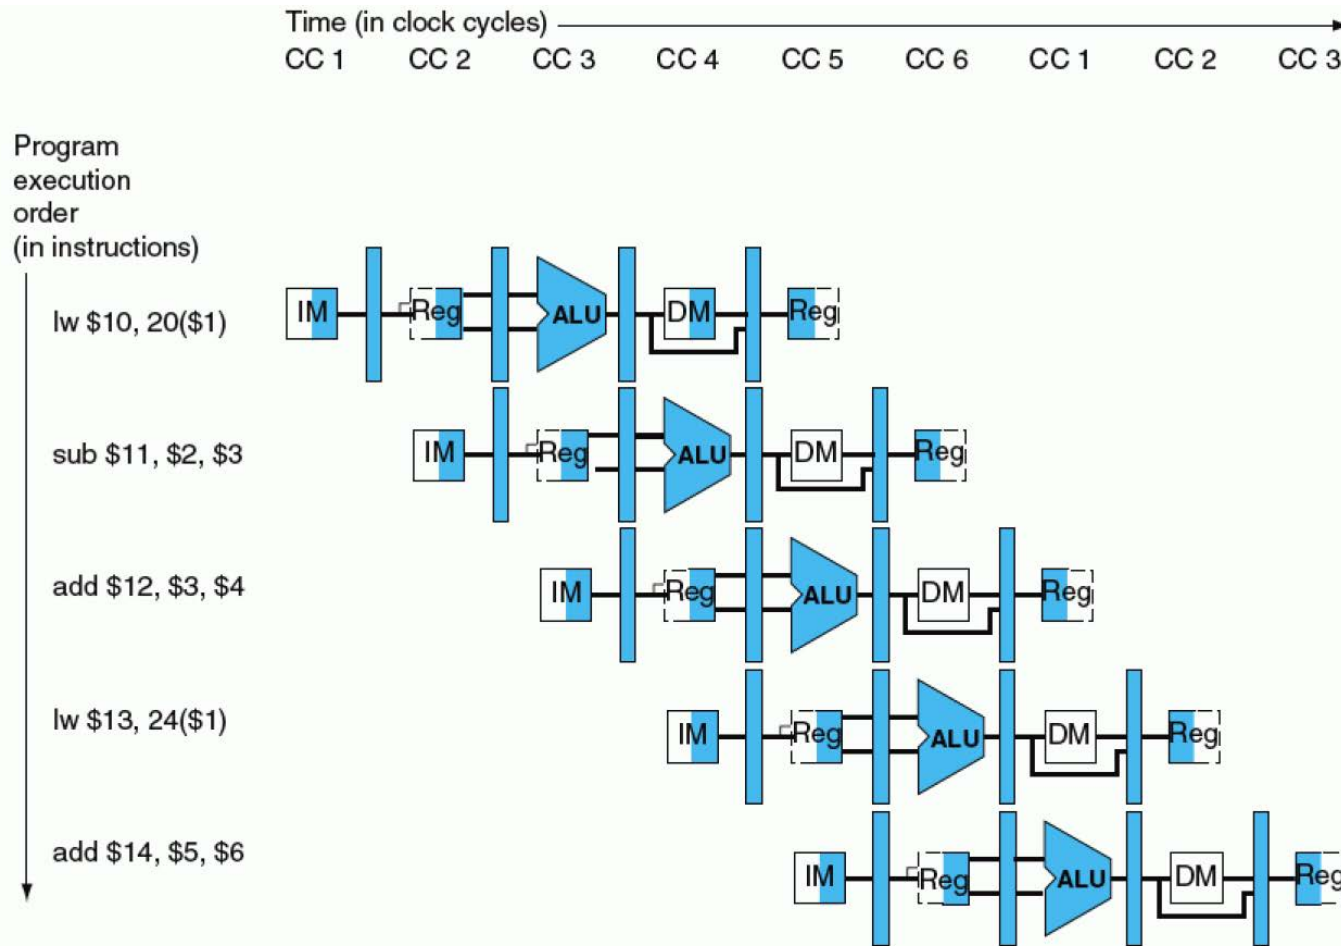
- ❑ At each cycle, one instruction is fetched and sent to the processor
- ❑ Ideally, after pipeline is fully filled, one instruction completes each cycle

Multi-clock-cycle pipeline diagram : traditional view ¹⁵

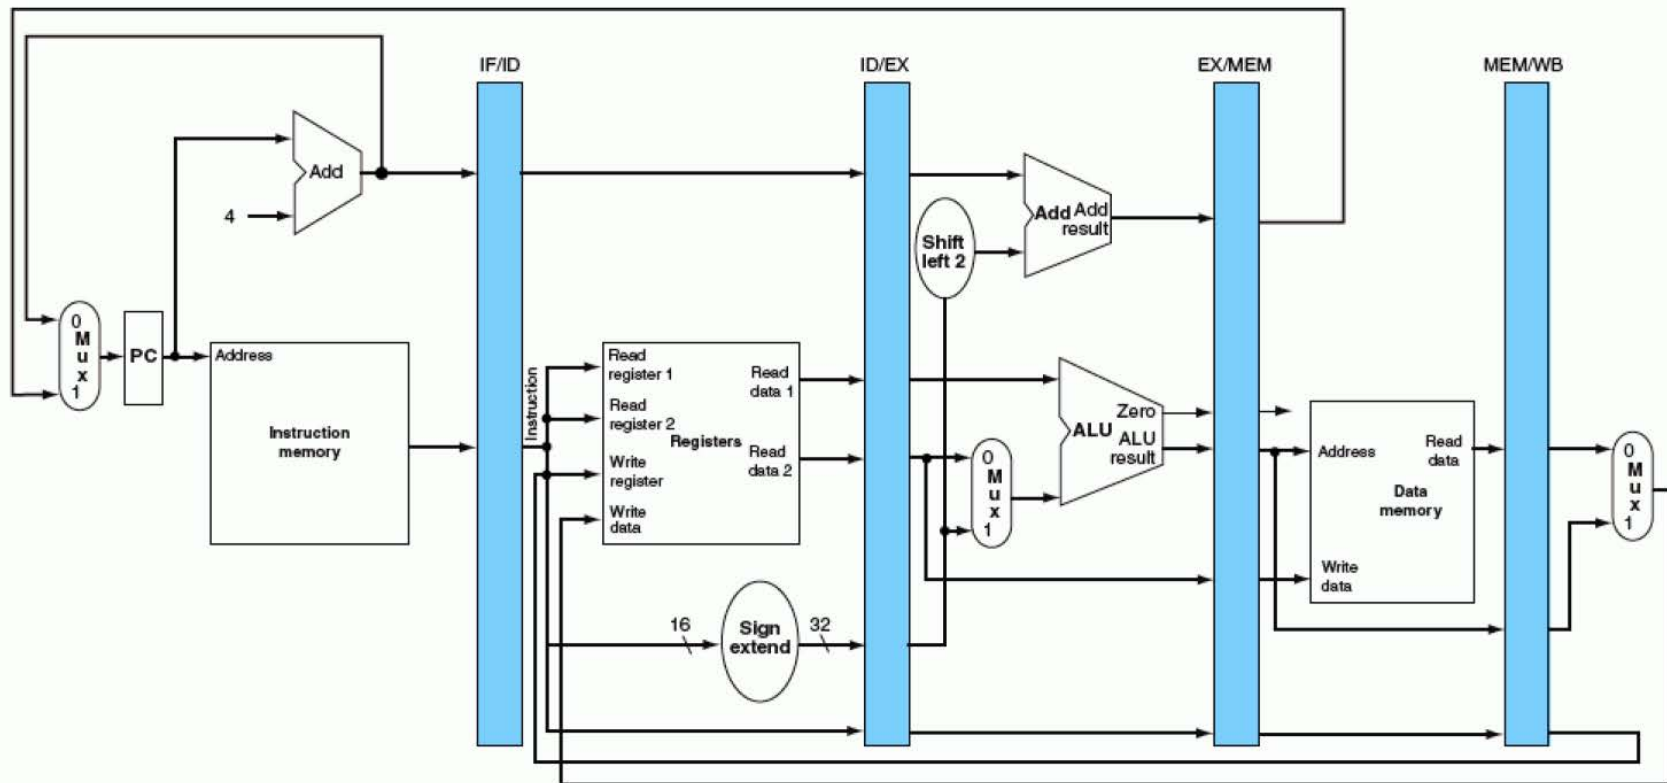
- ❑ The following diagram shows the execution of a series of instructions.



- ❑ The multi-clock-cycle form showing the hardware utilizations.



add \$14, \$5, \$6	lw \$13, 24 (\$1)	add \$12, \$3, \$4, \$11	sub \$11, \$2, \$3	lw\$10, 20(\$1)
Instruction fetch	Instruction decode	Execution	Memory	Write back



2- Dependencies and Hazards

Sometimes these dependences cause the pipeline to not fully fill

- ❑ Execution stops to wait for data or control to be produced
- i.e. next instruction cannot be executed in the next cycle

Data dependence

- ❑ Example:

```
lw      $1, 200($2)
add     $3, $4, $1
```
- ❑ **add** can't carry out **ID** until **\$1** has the updated value from the 5th stage of **lw**

Control dependence

- ❑ Example:

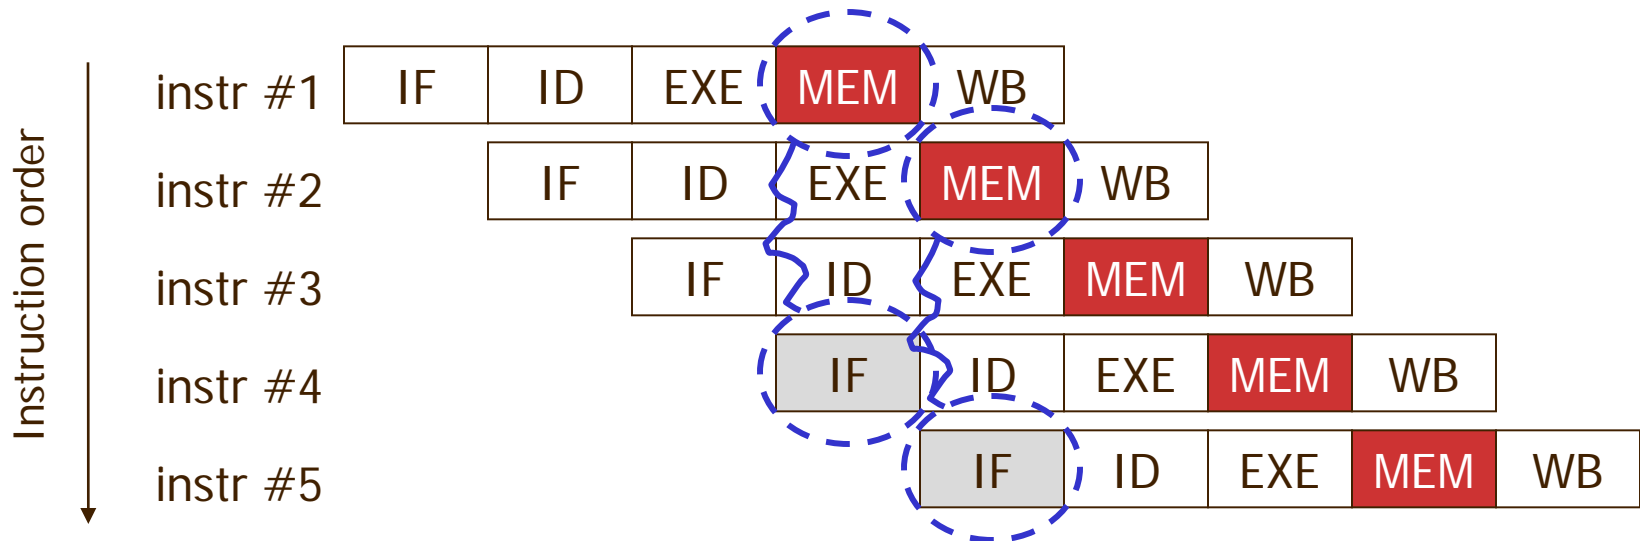
```
bne     $1, $2, target
add     $3, $4, $5
```
- ❑ **add** can't carry out **IF** until **bne** has completed the comparison in the 3rd stage of **bne**

- ❑ **Hazards** are situations in pipelining when the next instruction cannot be executed in the following clock cycle (or **pipeline stall**)
- Hazards reduce the performance of pipelining

Three types of pipelined hazards

- ❑ **Structural hazards:** hardware cannot support the combination of instructions to execute in the same clock cycle
- ❑ **Control hazards:** which instruction to execute next depends on the results of a previous instruction still in the pipeline
- ❑ **Data hazards:** an instruction depends on the results of a previous instruction still in the pipeline

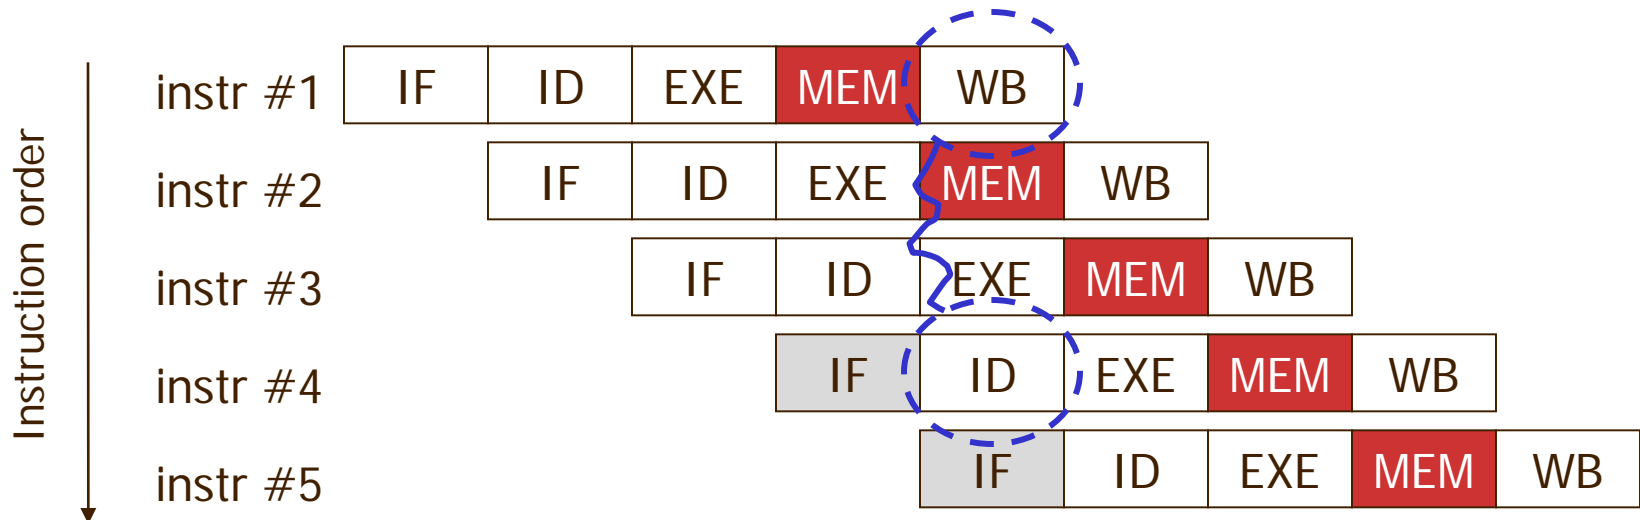
- ❑ If instructions #1 and #2 are load operations → structural hazard
- Instruction fetch (#4, #5) and data load (#1, #2) need memory access



Read same memory twice in same clock cycle
(assumption: memory can only service one at a time)

- ❑ Solution?
- Add memory ports to allow parallel accesses to independent addresses
- Use caching with separate Level 1 Data cache and Instruction cache

- ❑ If instr. #1 is a store operation → structural hazard with instr. #4
- As instr. #1 wants to write while instr. #4 wants to read the register file



Can't read and write to registers simultaneously

- ❑ Solution?

Fact: Register access is VERY fast;

- ❑ Takes less than half the time of ALU stage

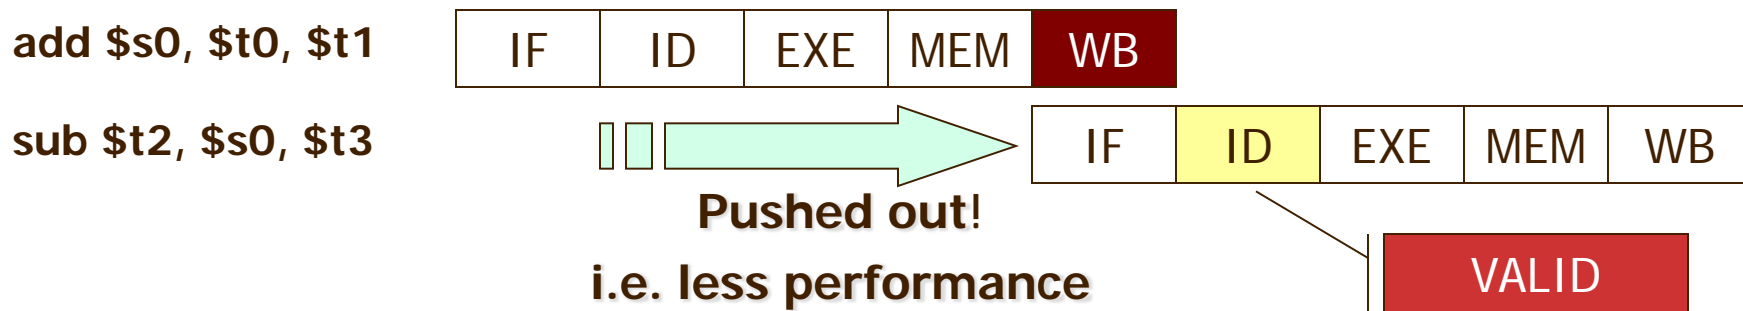
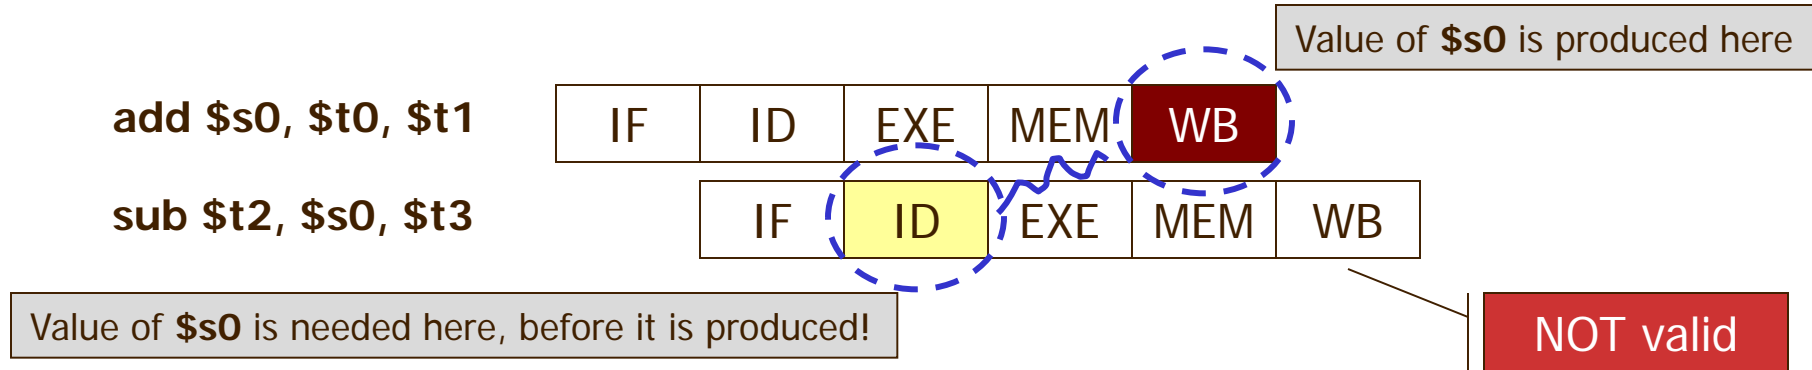
- ❑ Solution:

- ❑ always Write to registers during 1st half of each clock cycle
- ❑ always Read from Registers during 2nd half of each clock cycle
- ❑ Result: can perform write the Read during the same clock cycle

- ❑ Data hazards arise from the dependence of one instruction on an earlier one that is still in the pipeline

- ❑ e.g.

```
add $s0, $t0, $t1
sub $t2, $s0, $t3
```

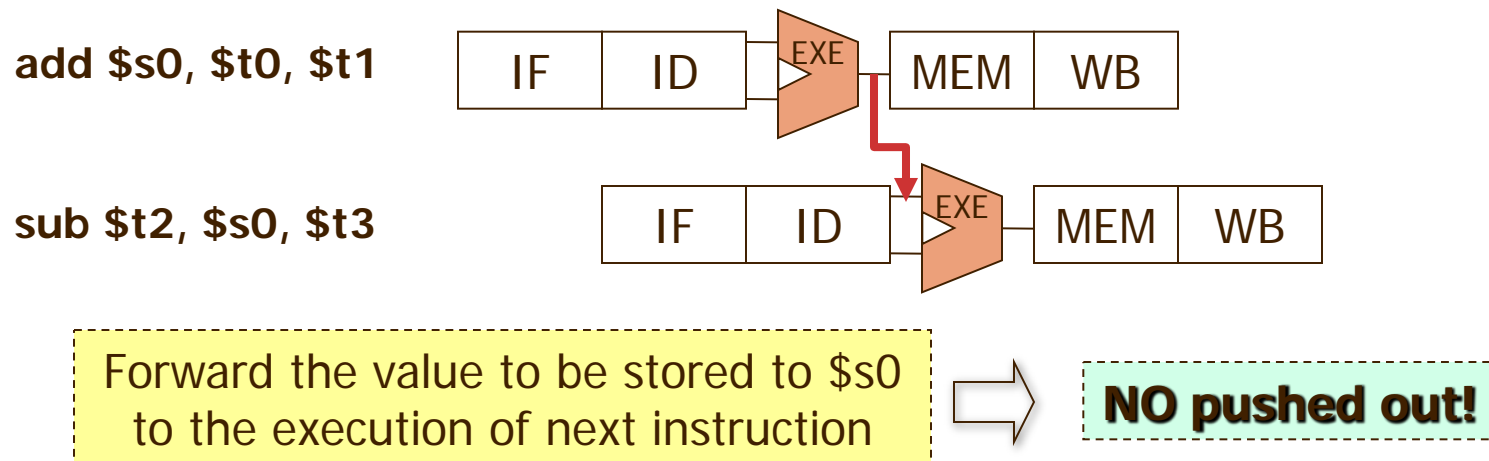


Observation:

- ❑ We don't need to wait for the instruction to complete before we try to resolve the data hazard

Solution:

- ❑ Add extra hardware to retrieve the missing item early from the internal resources, i.e. **forwarding** or **bypassing**



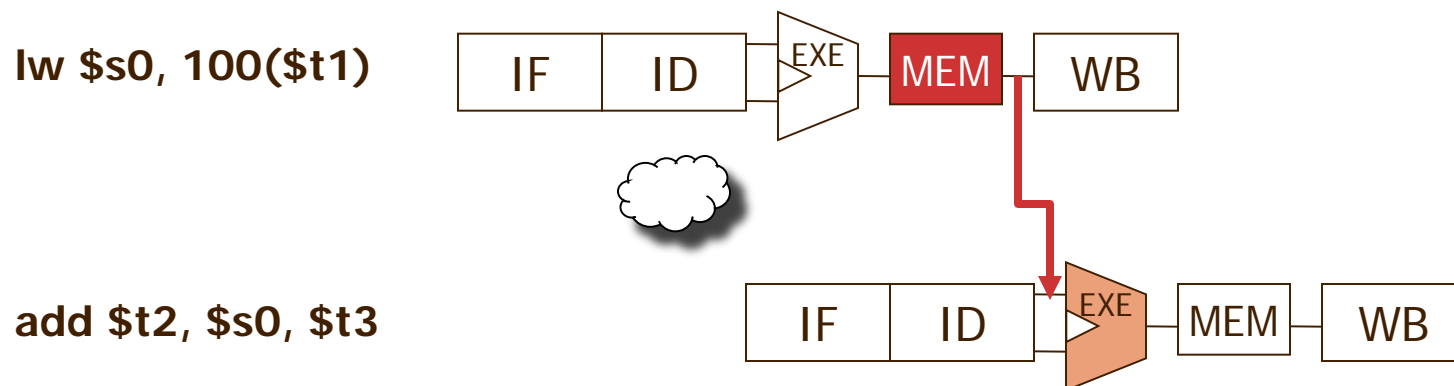
But, **pipeline stall** (nickname **bubble**) happens even with forwarding

- When a **R-format** instruction following a **load** tries to use the data

- e.g. `lw $s0, 100($t0)`
`add $t2, $s0, $t3`

- Why still bubble?

- Because the earliest time we can forward is from MEM, not EXE



We say **one** bubble is introduced, or processor has to stall for **one** cycle

- ❑ Consider code segment in C below
- ❑ All variables are in memory and are addressable as offsets from `$t0`

`A = B + E;`

`C = B + F;`

Original version:

```
lw      $t1, 0($t0)
lw      $t2, 4($t0)
add     $t3, $t1, $t2
sw      $t3, 12($t0)
lw      $t4, 8($t0)
add     $t5, $t1, $t4
sw      $t5, 16($t0)
```


Both `add` instructions have a hazard on the immediately proceeding `lw` instruction

Reordered version:


```
lw      $t1, 0($t0)
lw      $t2, 4($t0)
lw      $t4, 8($t0)
add     $t3, $t1, $t2
sw      $t3, 12($t0)
add     $t5, $t1, $t4
sw      $t5, 16($t0)
```

Reordering the instructions by filling the `bubbles` with other useful but independent “work”


lw \$t1, 0(\$t0)




lw \$t2, 4(\$t0)




add \$t3, \$t1, \$t2




sw \$t3, 12(\$t0)



lw \$t4, 8(\$t0)

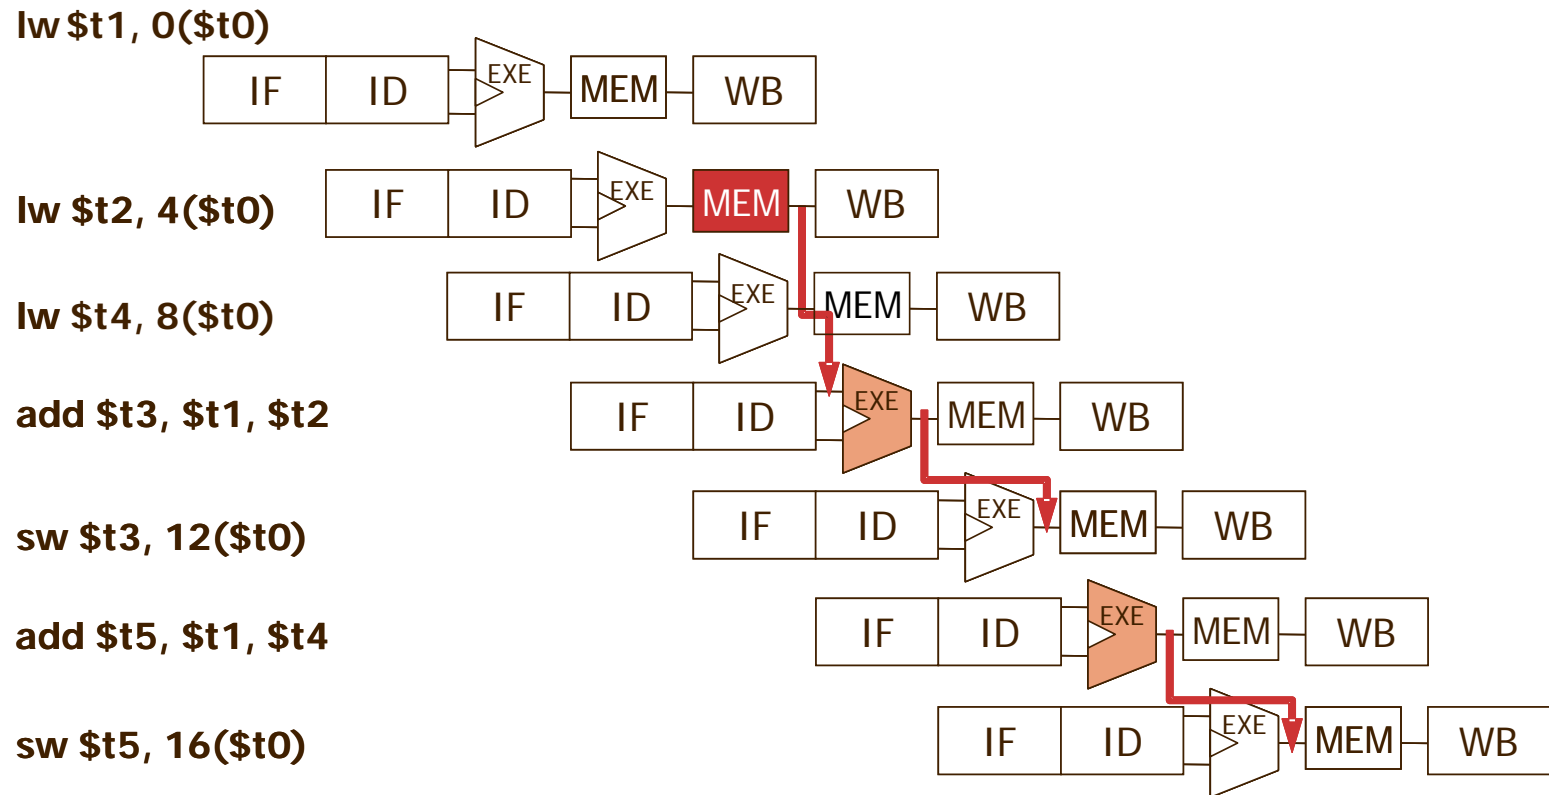


add \$t5, \$t1, \$t4




sw \$t5, 16(\$t0)





- Execution time = cycles for the ideal case + # of bubbles
= startup time + # of instructions + # of bubbles

Original version:



```
lw      $t1, 0($t0)
lw      $t2, 4($t0)
add     $t3, $t1, $t2
sw      $t3, 12($t0)
lw      $t4, 8($t0)
add     $t5, $t1, $t4
sw      $t5, 16($t0)
```

Reordered version:

```
lw      $t1, 0($t0)
lw      $t2, 4($t0)
lw      $t4, 8($t0)
add     $t3, $t1, $t2
sw      $t3, 12($t0)
add     $t5, $t1, $t4
sw      $t5, 16($t0)
```

- Execution time (original) = 4 + 7 + 2 = 13
- Execution time (reorder) = 4 + 7 = 11

- Assume `$s0` carry `A`, `$s1` carry `i`, `$s2` carry `n`
for (`i=2; i<n; i++`)
 `A[i] = A[i-1] + C; // C is a constant`

Assembly version:

```
addi    $t0, $s0, 0
LOOP:
addi    $t0, $t0, 4
lw      $t1, 0($t0)
addi    $t1, $t1, C
sw      $t1, 4($t0)
addi    $s1, $s1, 1
bne     $s1, $s2, LOOP
```

Assembly version:


```
addi    $t0, $s0, 0
LOOP:
addi    $t0, $t0, 4
lw      $t1, 0($t0)
addi    $s1, $s1, 1
addi    $t1, $t1, C
sw      $t1, 4($t0)
bne     $s1, $s2, LOOP
```

- Control hazards arise from the need to make a decision based on the results of one instruction while others are executing

e.g. `bne $1, $2, target`

`add $3, $4, $5`

`target: sub $3, $3, $6`



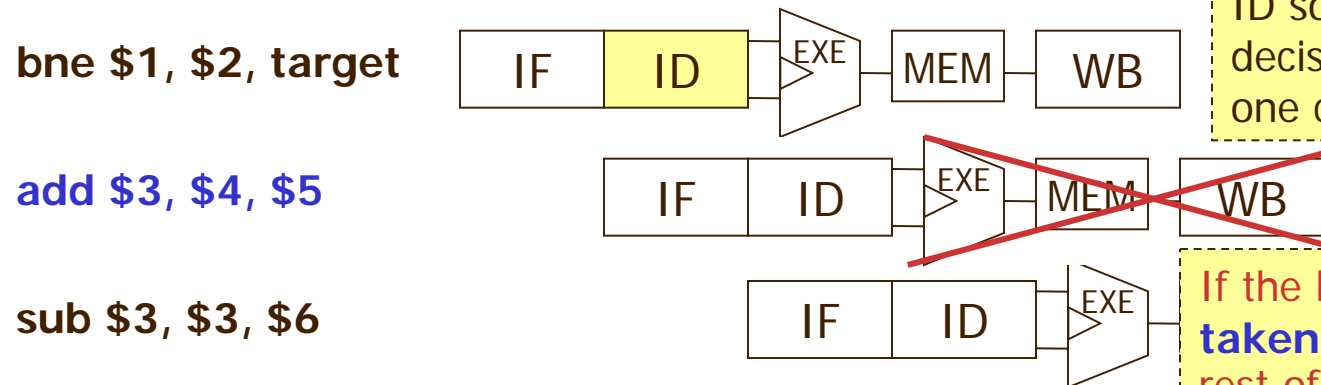
- Simple approach to reduce the stalling cycles from 2 to 1:
 - Dedicate a **separate comparator in ID stage** for branch comparison


```

        bne    $1, $2, target
        add    $3, $4, $5
target:  sub    $3, $3, $6

```

- ❑ Rather than stall, fill the bubble with work from **add**
- ❑ Basically, **predict branch not taken**. Why it works?



Reason:

- ❑ If the branch is really **not taken**, **add** already started, no pipeline stall!
- ❑ So, **reduce the frequency of stalling** due to branches

- ❑ Pipelining improves processor performance by overlapping instructions
 - ❑ **Pipelining principles**
 - ❑ The number of pipe stages decides the potential speedup
 - ❑ But, there is a limit in the number of pipe stage we can have

- ❑ Pipeline hazards
 - ❑ **Structural hazard**
 - ❑ **Data hazard**
 - ❑ **Control hazard**

- ❑ Techniques to reduce the negative impact of pipeline hazards
 - ❑ **Forwarding**
 - ❑ **Reorder the code**
 - ❑ **Make branch decision early**
 - ❑ **Fill bubble with potential work**